

Task 2: Values Audit

OCBC Bank Ltd - First-Person Values Audit for AI Governance in KYC Source-of-Wealth Review

Course: MH6822 Regulatory Technology

Selected entity: OCBC Bank Ltd

Jurisdictions: Singapore and United States

Domain: AI model risk management and AI governance for AML/KYC source-of-wealth due diligence

1. My company's mission statement

My hypothetical company is called Soteria RegTech.

Mission statement: Soteria RegTech helps banks use AI in high-risk compliance workflows without allowing automation to replace accountable human judgment.

The purpose of the company is not to help banks make compliance look neat on a dashboard. The purpose is to help banks test whether AI-assisted compliance decisions are supported by evidence, whether humans are genuinely reviewing AI outputs, and whether the correct jurisdictional controls are being applied.

This matters for OCBC because the selected use case is not a low-risk productivity tool. OCBC has a real Singapore headquarters and a real international network that includes the United States that have been mentioned in Task 1.

What my company is good at

Soteria RegTech is not a full AML transaction-monitoring vendor and not a law firm. It does not claim that software can make a bank MAS compliant or US compliant by itself. The company is good at four things:

Capability	What it means in this project
Jurisdiction-aware compliance logic	Translating different regulatory expectations into product rules without pretending that Singapore and the US are the same.
AI governance workflow design	Showing where AI can assist and where human

	judgment must remain accountable.
Evidence testing	Checking whether AI-generated source-of-wealth conclusions are actually supported by documents.
Auditability	Creating records that compliance, model risk, internal audit, and senior management can review.

Core value	Meaning
Accountability	AI can assist, but a named human or committee must remain responsible for final decisions.
Evidence over appearance	A polished AI-generated report is not enough if the underlying documents do not support it.
Jurisdictional honesty	The product must not hide real differences between Singapore and US expectations.
Risk integrity	The tool should detect genuine risk, not only help the bank produce better paperwork.
Commercial restraint	The company should not sell compliance comfort when the evidence is weak.

Company stage, size, and aspirations

I imagine Soteria RegTech as an early-stage Singapore-based RegTech company with around 20 employees.

Function	Approximate headcount
Data science and machine learning engineering	6
Software engineering and product	5
Compliance and regulatory specialists	4

Customer success and implementation	2
Sales and partnerships	2
Operations and finance	1

The company has a working prototype and is preparing for institutional pilots. Its first target market is Singapore-headquartered banks and foreign banks with Singapore operations. This is because MAS has already published AI governance materials for financial institutions, including its 2024 AI Model Risk Management information paper [3]. The second target market is banks with US branches, agencies, or wholesale operations, which makes OCBC a useful example because it is Singapore-headquartered but also has US banking agencies .

Period	Revenue aspiration	Geographic focus
Year 1	S\$500,000-S\$750,000	Singapore pilots
Year 2	S\$2 million-S\$4 million	Singapore and selected US-facing banks
Year 3	S\$8 million-S\$12 million	Singapore, Hong Kong, Australia, and US-facing banking clients

I have not done a full market-sizing study, so I would not claim that these figures are proven. They are assumptions based on the idea that banks are already experimenting with AI in compliance workflows and will need stronger governance tools as these systems become more embedded.

The commercial opportunity is not simply that banks need AI governance. The narrower opportunity is that banks are already deploying AI into KYC and source-of-wealth workflows, and the speed gain creates a governance problem that spreadsheets and policy memos cannot monitor continuously.

2. Whose perspective does my tool serve?

My tool primarily serves the bank's control function, especially the Chief Compliance Officer, Financial Crime Compliance team, Model Risk Management team, AI Governance Committee, and Internal Audit.

This is an important distinction. I am not designing the tool primarily for relationship managers or the front office, even though they may use some outputs. Relationship managers benefit from faster onboarding and faster report drafting, but they may also have incentives to reduce friction for valuable clients. The compliance function needs a tool that can slow down or challenge a file when the AI-generated report looks complete but the evidence remains weak.

Why I chose the compliance perspective

I chose the compliance and model-risk perspective for three reasons. First, the compliance function has the strongest reason to care about whether AI is creating real risk. Second, the bank has the ability to pay for and implement the tool. Third, my hypothetical company's competency is in governance and evidence testing, not in front-office sales enablement.

This choice is commercially practical, but it also creates a values risk. Since the bank pays for the tool, there may be pressure to make the tool produce reassuring outputs. A bank might prefer fewer escalations, faster onboarding, and dashboards that show that everything is under control. My product should resist that pressure by making uncertainty, weak evidence, contradictions, and human rubber-stamping visible.

Stakeholder tensions

There are real tensions between the stakeholders. The bank wants efficiency, consistency, and regulatory defensibility. This is clear from the Bank of Singapore source-of-wealth use case: Bank of Singapore says its Source of Wealth Assistant is an agentic AI-powered tool that can reduce source-of-wealth report preparation from around 10 days to one hour [2].

The front office benefits from this speed because faster source-of-wealth documentation can make onboarding smoother. However, the compliance function needs to know whether the AI-generated report is actually reliable. A faster report is not necessarily a better report.

The regulator wants effective AML/CFT controls, not just faster paperwork. In Singapore, MAS Notice 626 sets AML/CFT requirements for banks [7]. In the US, FinCEN states that the CDD Rule strengthens customer due diligence requirements for US banks and other covered financial institutions [9].

The customer wants a faster and less intrusive onboarding process. A legitimate customer may benefit if AI makes source-of-wealth review more consistent and less dependent on one relationship manager's drafting ability. However, the customer can also be harmed if the tool misreads documents,

exaggerates risk, causes unnecessary escalation, or exposes sensitive personal and financial information.

The public-good perspective is also important. In AML/KYC, the public harm is not abstract. The harm is that illicit wealth, corruption proceeds, sanctions-linked funds, or other financial crime proceeds may be made legitimate-looking through a respected financial institution. A false negative in source-of-wealth review is therefore not just an internal compliance error. It can weaken the integrity of the financial system.

My values position is that the tool should serve the bank's compliance function first, but it should not merely protect the bank from paperwork failure. It should also make it harder for weak evidence to be converted into false confidence.

3. Is my tool measuring genuine risk or documenting compliance?

My tool should measure genuine risk, not merely document compliance. The central problem in this domain is that an AI system can produce a source-of-wealth report that looks coherent, complete, and professional even when the underlying evidence is weak. A compliance team may then mistake narrative quality for evidentiary quality.

The tool should ask: Is the source-of-wealth conclusion actually supported by evidence, or has the AI merely produced a convincing summary?

Design choice I would include: Evidence Sufficiency and Contradiction Score

The main design choice I would include is an Evidence Sufficiency and Contradiction Score. This feature would test whether the AI-generated source-of-wealth narrative is actually supported by the documents and data in the file. It is included because it improves risk detection, even though a narrow compliance-documentation tool might not require it.

Metric	What it measures	Why it measures genuine risk
Evidence coverage ratio	Percentage of major wealth claims linked to supporting documents	Shows whether the file is substantively supported, not just completed.
Unsupported claim rate	Percentage of AI-generated claims with no evidence link	Detects hallucinated or overconfident AI narratives.
Contradiction count	Number of conflicts between	Detects unreliable source-of-

	the AI report and source documents	wealth conclusions.
Document consistency score	Whether names, dates, amounts, entities, and jurisdictions match across documents	Identifies weak or inconsistent evidence.
Human challenge rate	How often reviewers amend, reject, or escalate AI outputs	Tests whether human review is meaningful.
Override quality score	Whether approvals of weak files include specific reasons	Detects rubber-stamping.
Post-review escalation rate	How often AI-assisted approvals are later escalated	Detects workflow failure or model drift over time.

The tool should not allow a file to receive a low-risk governance status if the unsupported claim rate or contradiction count exceeds a defined threshold, even if all mandatory documentation fields have been completed.

This is the key difference between measuring real risk and documenting compliance. A documentation tool asks: Was the report completed? My tool should ask: Is the report actually supported?

Design choice I would exclude: a simple compliant badge

I would deliberately avoid a simple green label saying MAS compliant, US compliant, AI model approved, or KYC file passed. That feature would be commercially tempting because it gives senior management a simple answer. But it would serve documentation over substance. It could also encourage users to treat the tool as a substitute for judgment.

Instead, the tool should display specific findings such as unsupported claims, unresolved contradictions, missing documents, weak human override reasons, jurisdictional governance gaps, model drift indicators, and unresolved privacy or data-handling risks.

Uncomfortable design choice: red friction state

I would include a red friction state. If the AI-generated source-of-wealth report contains unsupported material claims or unresolved contradictions, the tool should block straight-through approval and

require senior compliance review. This is commercially inconvenient because it can slow onboarding. However, it reflects my values position that source-of-wealth review should not become only a fast-writing exercise.

The tool should not optimise primarily for reducing escalations. A lower escalation rate may indicate efficiency, but it may also indicate suppressed risk detection. Escalation reduction should only be treated as a positive outcome if it is paired with strong evidence-quality metrics, low contradiction rates, and acceptable post-review outcomes.

Jurisdiction-aware risk logic

The tool must also be jurisdiction-aware because Singapore and the US are not making identical regulatory choices.

In Singapore mode, the tool should align with MAS's AI governance materials. MAS's FEAT principles provide foundational principles for the responsible use of AI and data analytics in financial services [5]. MAS's Veritas initiative aims to help financial institutions evaluate AI and data analytics solutions against fairness, ethics, accountability, and transparency principles [6]. MAS's 2024 AI Model Risk Management information paper specifically concerns AI and Generative AI model risk management practices in financial institutions [3].

So, in Singapore mode, the tool should check AI governance ownership, human review, data quality, evidence sufficiency, fairness, accountability, transparency, model monitoring, and auditability.

In US mode, the tool should not falsely claim that agentic AI is fully covered by the latest US model-risk guidance. OCC Bulletin 2026-13 states that generative AI and agentic AI models are novel, rapidly evolving, and not within the scope of that guidance [10]. The Federal Reserve issued corresponding revised model risk guidance through SR 26-2 [1].

Therefore, in US mode, the tool should create a governance gap warning for generative or agentic AI. It should still require controls, but it should be honest that those controls come from broader AML, customer due diligence, internal governance, and supervisory risk-management expectations rather than from direct coverage under the 2026 model-risk guidance.

4. Who bears the cost if my tool gets it wrong?

The cost of error is not evenly distributed. Different stakeholders bear different harms depending on the failure mode.

Failure mode	Specific harm	Stakeholder most affected
False positive	A legitimate customer is incorrectly flagged as suspicious, causing delayed onboarding, intrusive document requests, or possible rejection.	Customer first; bank second.
False negative	A high-risk customer with illicit, unsupported, or misleading wealth is approved because the AI-generated report looked credible.	Public and regulator first; bank second.
Performance drift	The tool becomes less reliable as customer profiles, financial crime typologies, document formats, or economic conditions change.	Bank and regulator; customers may also be affected.
Jurisdictional misconfiguration	A Singapore file is assessed using US assumptions, or a US file is assessed using Singapore-only controls.	Bank compliance team first; regulator second.
AI hallucination	The AI invents facts, overstates evidence, or creates a persuasive but unsupported source-of-wealth narrative.	Bank and public first; customer if wrongly escalated.
Human over-reliance	Reviewers treat the AI draft as authoritative and stop challenging weak evidence.	Bank, regulator, and public.
Privacy or data-security failure	Sensitive source-of-wealth documents are exposed, retained improperly, or	Customer first; bank second.

	processed in the wrong environment.	
Threshold miscalibration	The tool is tuned to reduce workload instead of detecting true risk.	Public in false-negative cases; customers in false-positive cases.
Vendor design failure	My company designs the interface in a way that hides uncertainty or makes weak files look safe.	Bank, regulator, public, and customers.

The most serious error: false negative

In this domain, I consider a false negative the most serious error. A false negative means the tool helps the bank become comfortable with a customer whose source of wealth is actually suspicious, unsupported, or illicit. The public bears the first-order harm because AML/CFT controls exist to protect the financial system from financial crime, corruption proceeds, sanctions evasion, and terrorism financing. The bank then bears second-order harm through regulatory scrutiny, remediation cost, reputational damage, and loss of supervisory trust.

Because of this, I would accept a moderate increase in false positives for high-risk cases, provided the follow-up is respectful, explainable, and proportionate. This is not because customer friction is unimportant. It is because AML/KYC is designed to prevent harms that may not be visible to the individual customer at the point of onboarding.

The most visible customer harm: false positive

The most visible customer harm is a false positive. A legitimate customer may be asked for unnecessary documents, experience delays, or feel unfairly treated as suspicious. This is especially serious in private banking, where source-of-wealth review may involve personal financial history, family wealth, corporate ownership, tax records, property documents, and transaction information.

For that reason, the tool should not simply increase escalations without explanation. It should show why a file was escalated, what evidence is missing, and what specific issue needs human review.

The hardest error to detect: human over-reliance

The hardest failure to detect is human over-reliance. A bank may formally say that the system is human-in-the-loop, but in practice reviewers may rubber-stamp the AI-generated report because it is fast, polished, and consistent. This is why my tool should include a Human Challenge Log.

- whether the reviewer changed the AI-generated report;
- whether the reviewer challenged unsupported claims;
- whether contradictions were escalated;
- whether overrides were specific or generic;
- whether the same reviewer repeatedly approved weak files;
- whether senior compliance reviewed red-friction cases.

This matters because human-in-the-loop is not meaningful unless the human has enough time, information, authority, and incentive to disagree with the AI.

Privacy and data-security harms

Privacy is not a secondary issue in this use case. Source-of-wealth review can involve some of the most sensitive non-medical personal and financial data a customer provides to a bank. This may include income records, tax documents, family wealth information, corporate ownership records, property records, investment information, and transaction histories.

Therefore, the tool should not use customer documents to train general-purpose models. It should minimise retention of raw documents, separate evidence extraction from model-improvement pipelines, and require explicit governance approval before any sensitive data is reused for model development.

Vendor responsibility

Soteria RegTech also bears responsibility if the product design encourages false confidence. It is not enough for my company to say that the bank remains the final decision-maker. If the interface makes weak files look green, hides uncertainty, rewards lower escalation rates without measuring missed risk, or encourages reviewers to approve AI outputs without challenge, then the vendor has contributed to the failure.

For that reason, the tool should include uncertainty labels, evidence gaps, audit logs, human challenge indicators, and post-deployment monitoring. It should not only show workflow completion.

Task 2 conclusion

My values position is that the tool should serve the bank's compliance and model-risk functions, not the front office's desire for faster onboarding. The tool should make AI-assisted source-of-wealth review more evidentially reliable, not merely faster or easier to document. The hardest design problem is that the paying client benefits from speed and defensibility, while the public benefits from skepticism and escalation in weak-evidence cases.

For that reason, my tool should deliberately surface uncertainty, contradictions, unsupported claims, human rubber-stamping, model drift, privacy risks, and jurisdictional governance gaps, even when those findings slow down the business.

Central design principle: AI may assist compliance work, but it should not convert uncertainty into false confidence.

References

- [1] **Federal Reserve SR 26-2, Revised Guidance on Model Risk Management.** Provides Federal Reserve supervisory letter for revised model risk management guidance. <https://www.federalreserve.gov/supervisionreg/srletters/SR2602.htm>
- [2] **Bank of Singapore, Source of Wealth Assistant media release.** Describes the agentic AI source-of-wealth use case and human review controls. <https://www.bankofsingapore.com/media-releases/2025/bank-of-singapore-deploys-agentic-ai-tool-to-automate-writing-of-source-of-wealth-reports.html>
- [3] **MAS, Artificial Intelligence Model Risk Management information paper.** Sets out MAS good practices for AI and Generative AI model risk management. <https://www.mas.gov.sg/-/media/mas-media-library/publications/monographs-or-information-paper/imd/2024/information-paper-on-ai-risk-management-final.pdf>
- [4] **MAS Circular ID 18/24.** States that the AI MRM information paper was based on a thematic review of banks AI, including Generative AI, model risk management practices. https://www.mas.gov.sg/-/media/mas-media-library/regulation/circulars/id/id18_24/id18_24.pdf
- [5] **MAS, FEAT Principles.** Sets out fairness, ethics, accountability, and transparency principles for AI/data analytics in financial services. <https://www.mas.gov.sg/~media/MAS/News%20and%20Publications/Monographs%20and%20Information%20Papers/FEAT%20Principles%20Final.pdf>

- [6] **MAS, Veritas Initiative.** Describes MAS-industry work to operationalise responsible AI assessment against FEAT principles. <https://www.mas.gov.sg/schemes-and-initiatives/veritas>
- [7] **MAS, Notice 626.** Sets AML/CFT requirements for banks in Singapore. <https://www.mas.gov.sg/regulation/notices/notice-626>
- [9] **FinCEN, CDD Final Rule.** Describes US customer due diligence requirements for covered financial institutions. <https://www.fincen.gov/resources/statutes-and-regulations/cdd-final-rule>
- [10] **OCC Bulletin 2026-13, Model Risk Management: Revised Guidance.** States revised model risk management guidance and excludes generative AI and agentic AI from its scope. <https://www.occ.treas.gov/news-issuances/bulletins/2026/bulletin-2026-13.html>